



New York City Taxi Fare Prediction

<https://www.kaggle.com/c/new-york-city-taxi-fare-prediction>

[Click here](#) to view the notebook of the model we critique in this presentation

Team 12: Kate Luo, Ian Wesley McDaniel, David Tabert, Christopher Carlevato



Overview

- I. Objective
- II. Dataset
- III. Original Model
- IV. Our Model
- V. Feature Engineering
- VI. Results and Improvements



Objective

- Predict fare amount (inclusive of tolls) for a taxi ride in NYC given the pickup and dropoff locations
- Linear model → XGBoost
 - A decision tree algorithm, rather than a simple linear model
 - Beat linear model RMSE: \$5.74
 - Model was, on average, off by \$5.74
 - Average fare price: \$11.33
- Application: Provide riders an accurate cost estimate for their ride
 - Key user feature of ride hailing apps



Objective

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2}$$

- Square root of the squared difference between prediction and actual values (residuals) averaged over all of the sample
- Relative to MAE, RMSE penalizes large errors
 - Errors are squared before they are averaged
- Advantage
 - Model prediction error are in the units of the response variable



Dataset

- Separate CSV files for train & test data
 - Big data
 - 55M rows
- Read in 3M rows

```
NY_Train <- fread("train.csv", nrows = 3000000)
```

Features

- pickup_datetime - timestamp value indicating when the taxi ride started.
- pickup_longitude - float for longitude coordinate of where the taxi ride started.
- pickup_latitude - float for latitude coordinate of where the taxi ride started.
- dropoff_longitude - float for longitude coordinate of where the taxi ride ended.
- dropoff_latitude - float for latitude coordinate of where the taxi ride ended.
- passenger_count - integer indicating the number of passengers in the taxi ride.

Target

- fare_amount - float dollar amount of the cost of the taxi ride. This value is only in the training set; this is what you are predicting in the test set and it is required in your submission CSV.



NYC Taxi Stats

- NYC taxi [trend dashboard](#)
- Average 180 miles per shift
- Decline since launch of ride hailing apps in 2011



Original model

- Strengths:
 - Trained on ~10,000,000 samples
 - Simple
- Weakness:
 - Truth is very rarely linear
 - Large RMSE
 - Multicollinearity
 - Little feature engineering
 - Average lat and lon between pickup and dropoff
 - Data quality
 - Little data cleansing happened
 - Only NaNs removed from lat/lon columns
 - Garbage in → Garbage out



Our Model: XGBoost

- Looked at some published notebooks and found that XGBoost performed well with this data
 - [Mitchell O'brien's notebook](#)
- XGBoost advantages
 - Tree-based method which does not assume linearity
 - Slow and weak but effective learners
 - Include a validation stopping rule to avoid overfitting
 - More regularized model to control overfitting
 - Feature importance
- Disadvantage:
 - Computationally expensive



XGBoost Algorithm:

-
1. Set $\hat{f}(x) = 0$ and $r_i = y_i$ for all i in the training set.
 2. For $b = 1, 2, \dots, B$, repeat:
 - (a) Fit a tree \hat{f}^b with d splits ($d + 1$ terminal nodes) to the training data (X, r) .
 - (b) Update \hat{f} by adding in a shrunk version of the new tree:

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x). \quad (8.10)$$

- (c) Update the residuals,

$$r_i \leftarrow r_i - \lambda \hat{f}^b(x_i). \quad (8.11)$$

3. Output the boosted model,

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x). \quad (8.12)$$



XGBoost Model: Added features

- Distance
 - Haversine distance
- Time Features
 - Year
 - Month
 - Day of the week
 - Etc.
- Added distance from pickup location to popular destinations as features
 - Airports and major city landmarks
 - e.g. distance to JFK

Your most recent submission

Name	Submitted	Wait time	Execution time	Score
team12_sub.csv	just now	0 seconds	0 seconds	3.04896

Complete

XGBoost Results

- Learning Rate 0.05 \rightarrow 0.01
- 8-leaf trees \rightarrow 6-leaf trees
- 6,000 trees \rightarrow 500 trees
- Took out certain unimportant features and reran model
- Trained until validation RMSE had not improved in 10 rounds
- RMSE: 3.04896
- Outperformed simple linear model RMSE by 2.69
- Unsurprisingly, distance is the most important feature



Future improvement

- Use entire training dataset
 - Used 5% of all training data
- More feature engineering
- Hyperparameter tuning
- Rerun model with principal components
- Try other models like Random Forest to see if they outperform XGBoost



Sources

- <https://www.kaggle.com/obrienmitch94/nyc-taxi-fare-prediction>
 - Used this code as a prototype for our model, but improved model by changing hyperparameters and eliminating unnecessary features